

Assesment Report CLO-2 Case Based

K-Means Algorithm



Created by:

Muhammad Rizky Aulia Gobel – 1304211002

IF-45-02.1PJJ

Long-Distance Informatics major

Informatic faculty

Telkom University

Bandung

2023

Statement

I do this assignment honestly and follow the rule of conduct of academic ethics in case I do not follow the rule I am ready to get a 0 (zero) mark for my assignment.

A handwritten signature in black ink, appearing to read 'Rizky' with a stylized flourish at the end.

Muhammad Rizky Aulia Gobel

Import library

For this assignment, the author used several libraries as follows

```
1. Import library and file

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import random
import seaborn as sns
import openpyxl
```

Overview the dataset

To run an .xlsx file, the author uses Google Colab mount to mount to the author's drive.

Then, the author creates a variable called data that is used to read the water-treatment file that has been converted to an .xlsx file previously. Then, the author runs data.head() which is used to display the first 5 rows.

```
[2] from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

[3] data = pd.read_excel('/content/drive/My Drive/Machine Learning/CLD-2/water_treatment.xlsx', index_col=0)
data.head()
```

Date	Q-E	ZN-E	PH-E	ORO-E	DOQ-E	SS-E	SSV-E	SED-E	COND-E	PH-P	...	COND-S	RD-ORO-P	RD-SS-P	RD-SED-P	RD-ORO-S	RD-DOQ-S	RD-ORO-G	RD-DOQ-G	RD-SS-G	RD-SED-G
1990-01-01	41230.0	0.35	7.6	120.0	344.0	136.0	54.4	4.5	593	7.5	...	903.0	NaN	62.6	53.3	NaN	62.5	86.7	71.8	87.5	99.4
1990-01-02	37306.0	1.40	7.9	165.0	470.0	170.0	76.5	4.0	1365	7.9	...	1481.0	NaN	50.0	94.4	65.9	73.6	86.7	79.4	89.4	100.0
1990-01-03	34535.0	1.00	7.8	232.0	518.0	220.0	65.5	5.5	1617	7.9	...	1492.0	32.6	62.4	95.0	81.3	59.9	87.5	71.8	85.9	99.8
1990-01-04	32527.0	3.00	7.8	187.0	460.0	180.0	67.8	5.2	1832	7.9	...	1590.0	13.2	57.6	95.5	85.3	70.4	85.0	77.2	83.3	100.0
1990-01-07	27780.0	1.20	7.6	199.0	466.0	186.0	74.2	4.5	1220	7.5	...	1411.0	38.2	46.6	95.0	84.9	61.1	89.4	73.8	86.6	99.6

5 rows x 38 columns


```
[4] data.describe()
```

	Q-E	ZN-E	PH-E	ORO-E	DOQ-E	SS-E	SSV-E	SED-E	COND-E	PH-P	...	COND-S	RD-ORO-P	RD-SS-P	RD-SED-P	RD-ORO-S	RD-DOQ-S	RD-ORO-G	RD-DOQ-G	RD-SS-G	RD-SED-G
count	509.000000	524.000000	527.000000	504.000000	521.000000	526.000000	516.000000	502.000000	527.000000	527.000000	...	526.000000	465.000000	523.000000	500.000000	487.000000	501.000000	491.000000	502.000000	519.000000	496.000000
mean	37226.567780	2.359065	7.810057	188.714286	406.896273	227.444867	61.393217	4.593825	1478.620493	7.829981	...	1494.817490	39.085806	58.518738	90.554200	83.448049	67.817365	89.013646	77.856574	88.963391	99.086290
std	6635.998806	2.749509	0.246175	62.062831	120.362974	135.940589	12.415264	2.743947	394.898019	0.226925	...	387.906161	14.782976	12.802547	8.944776	8.884256	11.617612	7.827651	8.888932	8.221612	4.457165
min	19050.000000	0.100000	6.900000	31.000000	81.000000	98.000000	13.200000	0.400000	651.000000	7.300000	...	683.000000	0.600000	5.300000	7.700000	8.200000	1.400000	19.600000	19.200000	10.300000	36.400000
25%	32888.000000	0.900000	7.600000	146.750000	325.000000	178.000000	55.700000	3.100000	1261.000000	7.700000	...	1234.500000	29.300000	50.550000	88.600000	81.400000	63.100000	87.600000	74.000000	87.600000	95.300000
50%	35590.000000	1.500000	7.800000	182.500000	400.000000	196.000000	64.300000	4.500000	1406.000000	7.800000	...	1432.000000	39.600000	59.500000	93.300000	85.400000	69.900000	90.200000	79.150000	90.700000	99.700000
75%	41372.000000	3.000000	8.000000	223.000000	478.000000	242.000000	69.725000	5.500000	1672.000000	8.000000	...	1692.250000	49.200000	66.900000	96.000000	87.900000	75.300000	92.450000	83.400000	93.100000	100.000000
max	60081.000000	33.500000	8.700000	438.000000	941.000000	2008.000000	85.000000	36.500000	3230.000000	8.500000	...	3950.000000	79.100000	96.100000	100.000000	94.700000	96.800000	97.000000	98.100000	99.400000	100.000000

8 rows x 38 columns

Preprocessing

The author then ran data.describe() to check the contents of the dataset as a whole. What surprised the author was that the values in the count series had different values. There were some with 526 data points, 527 data points, 509 data points, and so on.



```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 527 entries, 1990-01-01 to 1991-10-30
Data columns (total 38 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Q-E         509 non-null    float64
1   ZN-E        524 non-null    float64
2   PH-E        527 non-null    float64
3   DBO-E       504 non-null    float64
4   DQO-E       521 non-null    float64
5   SS-E        526 non-null    float64
6   SSV-E       516 non-null    float64
7   SED-E       502 non-null    float64
8   COND-E      527 non-null    int64
9   PH-P        527 non-null    float64
10  DBO-P       487 non-null    float64
11  SS-P        527 non-null    int64
12  SSV-P       516 non-null    float64
13  SED-P       503 non-null    float64
14  COND-P      527 non-null    int64
15  PH-D        527 non-null    float64
16  DBO-D       499 non-null    float64
17  DQO-D       518 non-null    float64
18  SS-D        525 non-null    float64
19  SSV-D       514 non-null    float64
20  SED-D       502 non-null    float64
21  COND-D      527 non-null    int64
22  PH-S        526 non-null    float64
23  DBO-S       475 non-null    float64
24  DQO-S       508 non-null    float64
25  SS-S        512 non-null    float64
26  SSV-S       510 non-null    float64
27  SED-S       499 non-null    float64
28  COND-S      526 non-null    float64
29  RD-DBO-P    465 non-null    float64
30  RD-SS-P     523 non-null    float64
31  RD-SED-P    500 non-null    float64
32  RD-DBO-S    487 non-null    float64
33  RD-DQO-S    501 non-null    float64
34  RD-DBO-G    491 non-null    float64
35  RD-DQO-G    502 non-null    float64
36  RD-SS-G     519 non-null    float64
37  RD-SED-G    496 non-null    float64
dtypes: float64(34), int64(4)
memory usage: 160.6 KB
```

The author filled in the missing data using the bfill and ffill methods. The methods work as follows:

bfill fills in the missing data from top to bottom.

ffill fills in the missing data from bottom to top.

The results are then displayed as follows:

```
[ ] for i in range(1, 6):
    data.ffill(inplace=True)
    data.bfill(inplace=True)
data
```

Date	Q-E	ZN-E	PH-E	DBO-E	DQO-E	SS-E	SSV-E	SED-E	COND-E	PH-P	...	COND-S	RD-DBO-P	RD-SS-P	RD-SED-P	RD-DBO-S	RD-DQO-S	RD-DBO-G	RD-DQO-G	RD-SS-G	RD-SED-G
1990-01-01	41230.0	0.35	7.6	120.0	344.0	136.0	54.4	4.5	993	7.5	...	903.0	32.6	62.8	93.3	85.9	62.5	86.7	71.8	87.5	99.4
1990-01-02	37386.0	1.40	7.9	165.0	470.0	170.0	76.5	4.0	1365	7.9	...	1481.0	32.6	50.0	94.4	85.9	73.6	86.7	79.4	89.4	100.0
1990-01-03	34535.0	1.00	7.8	232.0	518.0	220.0	65.5	5.5	1617	7.9	...	1492.0	32.6	62.4	95.0	81.3	59.9	87.5	71.8	85.9	99.8
1990-01-04	32527.0	3.00	7.8	187.0	460.0	180.0	67.8	5.2	1832	7.9	...	1590.0	13.2	57.6	95.5	85.3	70.4	85.0	77.2	83.3	100.0
1990-01-07	27760.0	1.20	7.6	199.0	466.0	186.0	74.2	4.5	1220	7.5	...	1411.0	38.2	46.6	95.0	84.9	61.1	89.4	73.8	86.6	99.6
...
1991-10-25	35400.0	0.70	7.6	156.0	364.0	194.0	63.9	5.5	1680	7.6	...	1840.0	47.3	61.3	94.0	76.4	58.0	86.5	82.4	90.7	99.8
1991-10-26	30964.0	3.30	7.7	220.0	540.0	184.0	62.0	3.5	1445	7.7	...	1337.0	47.3	38.6	93.3	76.4	87.0	92.7	95.0	91.8	95.7
1991-10-27	35573.0	7.30	7.6	176.0	333.0	178.0	64.0	3.5	1627	7.7	...	1799.0	47.3	40.4	95.0	76.4	72.9	90.9	79.9	91.8	98.6
1991-10-29	29801.0	1.60	7.7	172.0	400.0	136.0	70.1	1.5	1402	7.7	...	1468.0	32.4	40.4	88.0	87.8	77.6	91.3	85.3	83.8	96.7
1991-10-30	31524.0	1.60	7.9	172.0	478.0	204.0	64.7	6.0	1798	7.9	...	1568.0	32.4	43.9	65.3	87.8	75.3	91.3	81.2	89.7	99.2

527 rows x 38 columns

```
[ ] data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 527 entries, 1990-01-01 to 1991-10-30
Data columns (total 38 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    Q-E         527 non-null    float64
1    ZN-E         527 non-null    float64
2    PH-E         527 non-null    float64
3    DBO-E         527 non-null    float64
4    DQO-E         527 non-null    float64
5    SS-E         527 non-null    float64
6    SSV-E         527 non-null    float64
7    SED-E         527 non-null    float64
8    COND-E        527 non-null    int64
9    PH-P         527 non-null    float64
10   DBO-P         527 non-null    float64
11   SS-P         527 non-null    int64
12   SSV-P         527 non-null    float64
13   SED-P         527 non-null    float64
14   COND-P        527 non-null    int64
15   PH-D         527 non-null    float64
16   DBO-D         527 non-null    float64
17   DQO-D         527 non-null    float64
18   SS-D         527 non-null    float64
19   SSV-D         527 non-null    float64
20   SED-D         527 non-null    float64
21   COND-D        527 non-null    int64
22   PH-S         527 non-null    float64
23   DBO-S         527 non-null    float64
24   DQO-S         527 non-null    float64
25   SS-S         527 non-null    float64
26   SSV-S         527 non-null    float64
27   SED-S         527 non-null    float64
28   COND-S        527 non-null    float64
29   RD-DBO-P      527 non-null    float64
30   RD-SS-P      527 non-null    float64
31   RD-SED-P      527 non-null    float64
32   RD-DBO-S      527 non-null    float64
33   RD-DQO-S      527 non-null    float64
34   RD-DBO-G      527 non-null    float64
35   RD-DQO-G      527 non-null    float64
36   RD-SS-G      527 non-null    float64
37   RD-SED-G      527 non-null    float64
dtypes: float64(34), int64(4)
memory usage: 160.6 KB
```

If you look at the first page, the rows 1990-01-01 and 1990-01-02 have NaN values. The NaN values are located in RD-DBO-P and RD-DBO-S, and on 1990-01-02, they are located in RD-DBO-P. After running bfill and ffill, the values are immediately filled in with the values below them. For example, the value of RD-DBO-S on row 1990-01-01 is filled in with the same value as RD-DBO-S on row 1990-01-02. In addition, with data.info(), the non-null values are

the same from Q-E to RD-SED-G, indicating that there are no more missing values and the dataset is ready to be used.

Now author can perform normalization, normalization will standardize the data so that it has a mean of 0 and a standard deviation of 1. This is important for many machine learning algorithms, as they often assume that the data is normally distributed, the method that is used is Z-Score, here's the code:

```
[ ] #Normalization
mean = data.mean()
std = data.std()

[ ] data = (data - mean)/std
data.head(5)
```

	Q-E	ZN-E	PH-E	DBO-E	DQO-E	SS-E	SSV-E	SED-E	COND-E	PH-P	...	COND-S	RD-DBO-P	RD-SS-P	RD-SED-P	RD-DBO-S	RD-DQO-S	RD-DBO-G	RD-DQO-G	RD-SS-G	RD-SED-G
Date																					
1990-01-01	0.623598	-0.729570	-0.853283	-1.099208	-0.523223	-0.672617	-0.578152	-0.025688	-1.229736	-1.454143	...	-1.526281	-0.432796	0.337649	0.313378	0.296026	-0.410102	-0.287504	-0.670416	-0.176696	0.069780
1990-01-02	0.035119	-0.346909	0.365362	-0.388432	0.528545	-0.422296	1.216078	-0.210630	-0.287721	0.308556	...	-0.034986	-0.432796	-0.665233	0.437137	0.296026	0.501085	-0.287504	0.169560	0.055064	0.207169
1990-01-03	-0.401341	-0.492685	-0.040853	0.669835	0.929218	-0.054177	0.323022	0.344195	0.350418	0.308556	...	-0.006604	-0.432796	0.306309	0.504642	-0.175352	-0.623533	-0.181150	-0.670416	-0.371852	0.161373
1990-01-04	-0.708746	0.236194	-0.040853	-0.040941	0.445071	-0.348672	0.509752	0.233230	0.894863	0.308556	...	0.246245	-1.756877	-0.069772	0.560896	0.234542	0.238401	-0.513506	-0.073891	-0.689006	0.207169
1990-01-07	-1.438528	-0.419797	-0.853283	0.148599	0.495155	-0.304498	1.029348	-0.025688	-0.654905	-1.454143	...	-0.215592	-0.050586	-0.931624	0.504642	0.193553	-0.525027	0.071441	-0.449370	-0.286476	0.115577


5 rows x 38 columns

After that, author will create the correlation matrix, correlation matrix is a table that summarizes the correlation between every pair of variables in a dataset. It is used to identify relationships between variables and can be used to inform feature selection and model building.

```
df = pd.DataFrame(data)
corr_matrix = df.corr()
corr_matrix
```

	Q-E	ZN-E	PH-E	DBO-E	DQO-E	SS-E	SSV-E	SED-E	COND-E	PH-P	...	COND-S	RD-DBO-P	RD-SS-P	RD-SED-P	RD-DBO-S	RD-DQO-S	RD-DBO-G	RD-DQO-G	RD-SS-G	RD-SED-G
Q-E	1.000000	0.036101	0.147464	-0.183635	-0.332340	-0.008384	-0.319980	-0.118319	-0.083469	0.163443	...	-0.072146	-0.086279	-0.024233	-0.052217	-0.041575	-0.040185	-0.074382	-0.160059	-0.018963	0.000374
ZN-E	0.036101	1.000000	-0.011092	0.005716	0.085096	0.091344	-0.098029	0.079847	0.052673	0.000461	...	0.069663	-0.066868	0.049984	0.016272	0.074306	0.042425	0.037761	0.058950	0.114517	0.031721
PH-E	0.147464	-0.011092	1.000000	0.224903	0.177341	-0.045714	0.184501	0.050052	0.272851	0.029089	...	0.233733	-0.039856	-0.153877	-0.150312	0.192597	0.219101	0.123425	0.141663	0.065359	0.020646
DBO-E	-0.183635	0.005716	0.224903	1.000000	0.493315	0.119519	0.230595	0.244184	0.203074	0.216045	...	0.197187	0.151443	-0.018956	-0.068779	0.173286	0.080904	0.283293	0.152571	0.018739	-0.034901
DQO-E	-0.332340	0.085096	0.177341	0.493315	1.000000	0.286192	0.249087	0.430992	0.308135	0.171856	...	0.275889	0.077419	0.125268	0.052885	0.103501	0.087526	0.096617	0.322210	0.111613	0.002688
SS-E	-0.008384	0.091344	-0.045714	0.119519	0.286192	1.000000	-0.542931	0.576708	0.083472	-0.100075	...	0.078518	0.179400	0.479287	0.140806	0.001950	-0.049955	0.059500	0.151772	0.258498	-0.002612
SSV-E	-0.319980	-0.098029	0.184501	0.230595	0.249087	-0.542931	1.000000	-0.140909	0.170660	0.219147	...	0.163223	-0.151194	-0.393917	-0.074077	0.132080	0.172460	0.048171	0.082931	-0.187719	0.025222
SED-E	-0.118319	0.079847	0.050052	0.244184	0.430992	0.576708	-0.140909	1.000000	0.191531	0.045189	...	0.160720	0.073082	0.318359	0.145227	0.077493	0.099712	0.083260	0.230171	0.198893	0.053105
COND-E	-0.083469	0.052673	0.272851	0.203074	0.308135	0.083472	0.170660	0.191531	1.000000	0.244974	...	0.901653	-0.056406	0.058835	0.019822	0.130875	0.101588	0.089655	0.067918	0.033439	0.052064
PH-P	0.163443	0.000461	0.029089	0.216045	0.171856	-0.100075	0.219147	0.045189	0.244974	1.000000	...	0.207641	-0.049267	-0.193449	-0.174160	0.154972	0.215552	0.093491	0.117419	0.032735	0.016392
DBO-P	-0.211799	0.035891	0.206113	0.665437	0.531477	0.128423	0.228529	0.259865	0.199735	0.191551	...	0.186242	0.433792	0.057738	0.047927	0.230488	0.121443	0.207350	0.110769	0.053024	-0.013676
SS-P	-0.024389	0.032169	-0.030850	0.119738	0.236597	0.877179	-0.503918	0.576002	0.091239	-0.092577	...	0.089035	0.237187	0.596318	0.201080	0.015797	-0.041909	0.056252	0.135114	0.249745	0.027093
SSV-P	-0.304727	-0.067507	0.187404	0.181835	0.214773	-0.500272	0.869826	-0.143197	0.159204	0.193047	...	0.158665	-0.194995	-0.486699	-0.129437	0.112906	0.136408	0.026243	0.031721	-0.202277	0.024236
SED-P	-0.127232	0.068273	0.056639	0.220626	0.382390	0.564995	-0.160176	0.841960	0.177300	0.028464	...	0.155806	0.136578	0.432917	0.237685	0.042036	0.054154	0.031776	0.160848	0.161033	0.026657
COND-P	-0.072753	0.050539	0.284004	0.218223	0.312388	0.080101	0.168250	0.186861	0.973147	0.258789	...	0.913724	-0.058843	0.048434	-0.001691	0.140565	0.108326	0.099731	0.059845	0.025279	0.055682
PH-D	0.136995	0.030728	0.822078	0.219710	0.186155	-0.071414	0.189635	0.074101	0.230338	0.894522	...	0.195721	0.010877	-0.133451	-0.155808	0.204378	0.245459	0.167333	0.170591	0.125775	0.090252
DBO-D	-0.168504	0.109348	0.287791	0.566562	0.521532	0.007623	0.374084	0.228408	0.301580	0.267760	...	0.310301	-0.299434	-0.130439	-0.080251	0.271291	0.189411	0.052571	0.074881	-0.104075	-0.065118
DQO-D	-0.214247	0.081829	0.263446	0.423646	0.642289	-0.024440	0.428550	0.221130	0.383584	0.280823	...	0.378793	-0.232138	-0.199775	-0.157249	0.105457	0.219868	-0.041241	0.003701	-0.116635	-0.075406
SS-D	-0.068053	0.096396	0.138566	0.304649	0.335086	0.204737	-0.084866	0.187115	0.115932	0.134976	...	0.128926	0.029164	-0.291469	-0.362509	0.015515	0.127654	0.019789	0.065174	0.047083	-0.045060
SSV-D	-0.332893	-0.026540	0.126125	0.147662	0.159855	-0.270203	0.682360	0.019642	0.125176	0.133208	...	0.118748	-0.169681	-0.092327	0.112661	0.123051	0.100968	0.036803	0.055168	-0.124632	0.038046
SED-D	-0.082840	0.064518	0.214605	0.246886	0.255927	0.043761	0.107328	0.152156	0.136324	0.225678	...	0.152943	-0.074668	-0.339577	-0.777487	0.003581	0.155584	0.000877	0.062964	0.012318	-0.050929

And here's the heatmap:




Algorithm(K-Means)

Before performing k-means clustering, author will choose 2 column from the dataset, the column that author pick is DQO-E and COND-E, and then DQO-E and COND-E will be converted into numPy array, then visualized with scatterplot

```
[ ] x = data['DQO-E']
    y = data['COND-E']
```

```
[ ] #Convert x and y into numpy array
    X = np.array(list(zip(x, y)))
```

```
plt.scatter(x, y, alpha=0.35)
plt.xlabel('DQO-E')
plt.ylabel('COND-E')
plt.show()
```



Then onto the K-Means, here's the code that author create:

```
def kmeans(X, k, max_iters=100):
    centroids = X[np.random.choice(range(len(X)), k, replace=False)]

    for _ in range(max_iters):
        distances = np.sqrt(((X - centroids[:, np.newaxis])**2).sum(axis=2))
        labels = np.argmin(distances, axis=0)
        new_centroids = np.array([X[labels == i].mean(axis=0) for i in range(k)])

        if np.all(centroids == new_centroids):
            break
        centroids = new_centroids

    return labels, centroids
```

def kmeans(X, k, max_iters=100) = function, X is the data points to cluster and represented numPy array, k is clusters to identify, and max_iters= is the maximum number to run the clustering algorithm

centroids = X[np.random.choice(range(len(X)), k, replace=False)] = initial centroid selection


```

for _ in range(max_iters):
    distances = np.sqrt(((X - centroids[:, np.newaxis])**2).sum(axis=2))
    labels = np.argmin(distances, axis=0)
    new_centroids = np.array([X[labels == i].mean(axis=0) for i in range(k)])

    if np.all(centroids == new_centroids):
        break

    centroids = new_centroids

```

= The clustering algorithm, it iterates until convergence or the maximum number of iterations is reached, each iteration involves the following steps:

1. Calculate distances;
2. Assign data points;
3. Recompute centroids, and
4. Check for convergence

then ended with

```
return labels, centroids
```


Elbow method

Elbow method is a method used in clustering to determine the optimum number of clusters in dataset, It is a heuristic method, doesn't guarantee that the optimum number will be found but it's relatively simple and effective method to quickly get the idea of the optimal numbers of clusters.

```
[ ] #Calculates the within-cluster sum of squares
def calculate_wcss(X, labels, centroids):
    wcss = 0
    for i, centroid in enumerate(centroids):
        cluster_points = X[labels == i]
        distance = np.sum((cluster_points - centroid)**2)
        wcss += distance
    return wcss
```

```
[ ] wcss_values = []
k_values = range(1, 11)
for k in k_values:
    labels, centroids = kmeans(X, k)
    wcss = calculate_wcss(X, labels, centroids)
    wcss_values.append(wcss)
```

```
plt.plot(k_values, wcss_values, 'bo-')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.title('Elbow Method')
plt.show()
```



`def calculate_wcss(X, labels, centroids)` = Define the function to calculate wcss, takes 3 arguments, X = data points to cluster, labels = an array of labels indicating the cluster assignment for each data point, and centroids = the final centroids representing centers of the cluster

for i, centroid in enumerate(centroids):

 cluster_points = X[labels == i]

 distance = np.sum((cluster_points - centroid)**2)

 wcss += distance

```
return wcss
```

= Calculation within wcss, a measure of how tightly the data points within each cluster are grouped together

```
wcss_values = [] = Create empty list
```

```
k_values = range(1, 11) = define range of k values to test
```

```
for k in k_values:
```

```
    labels, centroids = kmeans(X, k) = run the k-means function to obtain the cluster labels and centroids to the current k value
```

```
    wcss = calculate_wcss(X, labels, centroids) = calculate the wcss for the resulting clusters using the calculate_wcss() function
```

```
    wcss_values.append(wcss) = append the calculated wcss to the wcss_values list
```

= Calculating WCSS for Different k Values

Author decide that third cluster is the optimum, because it's resembles human-elbow, and for the silhouette method, here's the code:

Silhouette method

Silhouette method is a method used in cluster analysis to assess the quality of clustering algorithms. It measures how well each data point is assigned to its cluster based on two parameters that is:

-Cohesion

-Separation

```
[ ] #Calculates the Silhouette Coefficient for a given clustering result
def calculate_silhouette(X, labels):
    n = len(X)
    silhouette_values = np.zeros(n)

    for i in range(n):
        a = np.mean(np.sqrt(np.sum((X - X[i])**2, axis=1))[labels == labels[i]])

        b_values = []
        for j in range(k):
            if j != labels[i]:
                b = np.mean(np.sqrt(np.sum((X - X[i])**2, axis=1))[labels == j])
                b_values.append(b)

        b = np.min(b_values) if len(b_values) > 0 else np.inf

        silhouette_value = (b - a) / max(a, b)
        silhouette_values[i] = silhouette_value

    return np.mean(silhouette_values)

silhouette_values = []
k_values = range(2, 11)
for k in k_values:
    labels, centroids = kmeans(X, k)
    silhouette = calculate_silhouette(X, labels)
    silhouette_values.append(silhouette)

[ ] plt.figure(figsize=(8, 4))
plt.plot(k_values, silhouette_values, marker='o')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Silhouette Score')
plt.title('Silhouette scores for different number of clusters')
plt.xticks(k_values)
plt.show()
```

def calculate_silhouette(X, labels) = define the function for calculating silhouette

silhouette_values = initialize empty list

for i in range(n):

 a = np.mean(np.sqrt(np.sum((X - X[i])**2, axis=1))[labels == labels[i]])

= Calculate the cohesion by averaging the distances between current data point and other data points in the same cluster

 b_values = []

 for j in range(k):

 if j != labels[i]:

 b = np.mean(np.sqrt(np.sum((X - X[i])**2, axis=1))[labels == j])

 b_values.append(b)

 b = np.min(b_values) if len(b_values) > 0 else np.inf

= Calculate the separation by finding the minimum average distance between current data point and data points in all other clusters

$\text{silhouette_value} = (b - a) / \max(a, b)$

$\text{silhouette_values}[i] = \text{silhouette_value}$

= calculate the silhouette coefficient using the formula, then store the result

$\text{silhouette_values}[i]$

$\text{silhouette_values} = []$ = Create empty list

$k_values = \text{range}(2, 11)$ = define range of k values to test


for k in k_values:

 labels, centroids = kmeans(X, k) = run the k-means to obtain cluster labels for the current value(k)

 silhouette = calculate_silhouette(X, labels) = calculate average silhouette coefficient for the resulting clusters using the calculate_silhouette function

 silhouette_values.append(silhouette) = append the calculated average silhouette coefficient to the silhouette_value s

= Evaluating Silhouette for Different Cluster Numbers



Because author is curious what's the score, then author create this code

```
[ ] for i, k in enumerate(k_values):
    labels, centroids = kmeans(X, k)
    silhouette = calculate_silhouette(X, labels)
    print("Silhouette score for cluster =", k, "is", "{:.5f}".format(silhouette))

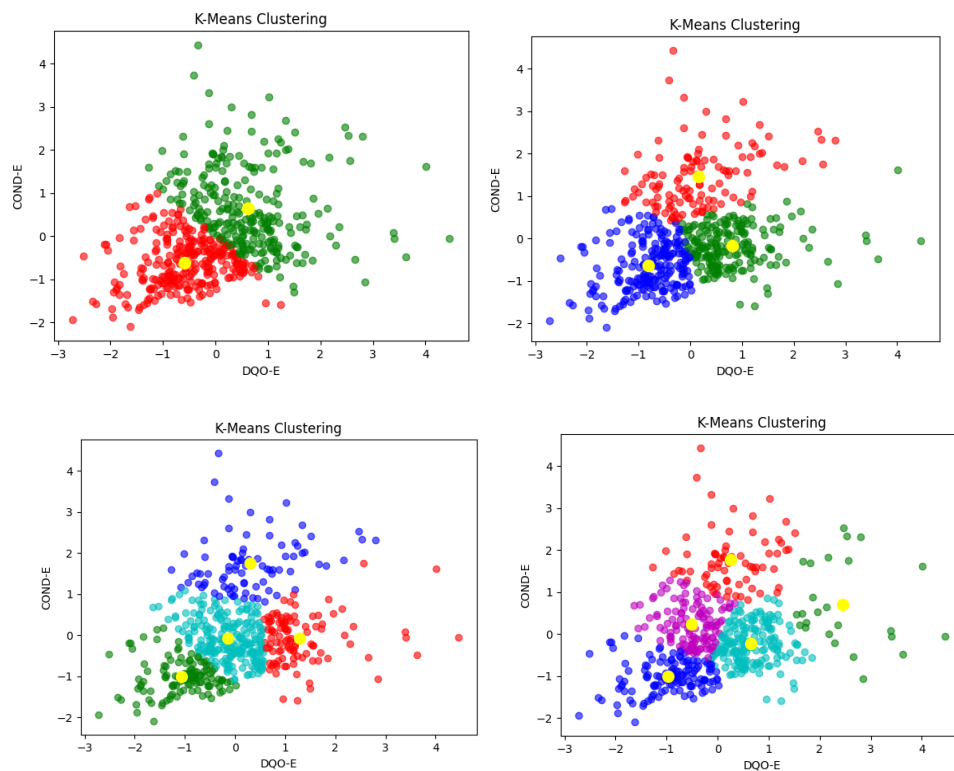
Silhouette score for cluster = 2 is 0.34679
Silhouette score for cluster = 3 is 0.38657
Silhouette score for cluster = 4 is 0.36955
Silhouette score for cluster = 5 is 0.36856
Silhouette score for cluster = 6 is 0.32595
Silhouette score for cluster = 7 is 0.33605
Silhouette score for cluster = 8 is 0.33328
Silhouette score for cluster = 9 is 0.34167
Silhouette score for cluster = 10 is 0.35039
```

With this, we know that third cluster is the most optimum cluster for silhouette method

Author conclude that elbow method and silhouette method has the same cluster which is 3.

Evaluate the results

Author create a test from second cluster to fifth cluster, and here's the result:



From left up -> right up -> bottom left -> bottom right : second cluster -> third cluster -> forth cluster -> fifth cluster.

Author conclude that the second cluster is the closest to 0.

During working for this assignment, troubles that author have met:

- Some of the column won't work really well with elbows method as they sometime spiked or drop randomly
- Creating kmeans without sklearn
- Choosing what pre-processing that suitable for this case

LINKS

Video Presentation: <https://www.youtube.com/watch?v=PbY9aTiKtZU>

GDrive: https://drive.google.com/drive/folders/1lOa5vJnLr0by2lKfXpB_e6st4GBIU6Ph?usp=sharing